



Five matched seeds • 4,000 REINFORCE iterations • bars $\pm$ sample SD

A. Learned alignment



B. Learned financial objective



C. Optimisation error in own reward

